

DH²A-KT: Dynamic Heterogeneous Hypergraph and Agentic Knowledge Tracing with Leakage-Controlled Graph Construction

Tuan Dao Minh, Khanh-Trinh Nguyen, Duong Nguyen Tien, Quoc Khanh Ngo, Van-Hau Nguyen, and Le Hoang Son

Abstract—Graph-augmented knowledge tracing (KT) often lacks leakage audits of its auxiliary graphs and rarely exposes frozen predictions to a gated explanation layer. We present DH²A-KT, a two-tier system that is tractable on a single consumer GPU. Tier 1 (DH²-KT) reuses the leakage-controlled, train-only concept-graph protocol of companion paper P0, groups audited prerequisite edges into *concept-prerequisite chain hyperedges* (the quantitative core evaluated here), and extends P0’s diagnostics with hyperedge-level leakage checks. Discussion-thread and teacher-intervention types remain data-blocked; session+Hint hyperedges and a limited observational hint ATE are demonstrated on FoundationalASSIST (not on XES3G5M). Tier 2 (AgentDG-KT) is a small-scale pilot: Diagnostician, Critic, and Tutor agents read Tier 1’s *frozen* outputs. On XES3G5M, DH²-KT v2 (chain hyperedges, graph-only interaction) achieves mean test AUC 0.752 versus reused GKT 0.834 under matched budget, while passing a hypergraph manipulation check ($\Delta\text{AUC} = 0.038$ at $p=0.9$ node-drop); P0 GKT also passes destruction at the same anchor with a larger mean drop, so reliance is not unique to DH²-KT. Junyi is used only for structural ground-truth overlap (not full model training). A 500-sample Tier 2 pilot with Qwen2.5-7B (Ollama) yields Critic flag rates of 12.4% under a strict prompt and 0% after probability-format calibration; a 40-sample dual-rater study rates faithfulness 4.01 ± 1.07 and usefulness 3.53 ± 0.71 (Likert 1–5). FoundationalASSIST supplies session+Hint hyperedges and a limited observational hint ATE (Section VII-E).

Index Terms—Knowledge tracing, hypergraph neural networks, heterogeneous graphs, large language model agents, causal inference, data leakage, educational data mining.

I. INTRODUCTION

Knowledge tracing (KT) estimates a learner’s latent mastery of knowledge components (KCs) from interaction history, supporting adaptive practice, intelligent tutoring, and curriculum recommendation. The field has progressed from probabilistic and sequence models [1] to graph-augmented models that propagate information over auxiliary concept graphs [2], [3], and, most recently, toward hypergraph representations that capture higher-order relations among knowledge components [4]–[6] and dynamic, temporally-evolving graph structure [7].

T. Dao Minh, K.-T. Nguyen, D. N. Tien, and V.-H. Nguyen are with the Department of Quality Assurance and Testing, Hung Yen University of Technology and Education, Hung Yen, Vietnam (e-mail: tuanymc@utehy.edu.vn; trinhnk@utehy.edu.vn; duongnt@utehy.edu.vn; nvhau66@gmail.com).

Q. K. Ngo and L. H. Son are with the Artificial Intelligence Research Center, VNU Information Technology Institute, Vietnam National University, Hanoi, Vietnam (e-mail: quockhanngo.official@gmail.com; sonlh@vnu.edu.vn). L. H. Son is the corresponding author.

This paper builds on – rather than replacing – [8]’s leakage-control protocol. Empirically, Tier 1 evaluates *concept-prerequisite* chain hyperedges under a graph-only encoder and a manipulation check. Discussion-thread and teacher-intervention hyperedges remain data-blocked; only session+Hint hyperedges and a limited observational hint ATE are demonstrated on FoundationalASSIST (weak propensity overlap; Sections IV-A–IV-D). Section IV-B extends P0’s diagnostics to the hyperedge setting. Tier 2 (Section V) is a deliberately-scoped pilot showing that an LLM-agent layer can read Tier 1’s frozen output without reopening the leakage question. On XES3G5M, graph-only interaction trades ≈ 8.2 AUC points versus GKT for a passing manipulation check; Tier 2 faithfulness at pilot scale is sensitive to Critic prompt design but reaches 0% flag rate after calibration (Section VI-C). Fig. 1 summarizes the two-tier pipeline.

A. Motivation

Three concrete gaps motivate this work:

- **Higher-order structure is under-audited in graph-KT.** Hypergraph-KT models [4]–[6] group multiple knowledge components into a single hyperedge, yet published systems rarely audit whether those hyperedges were constructed without leakage, and few expose an interface for entities beyond KCs (hints, teachers, forums, video) when such logs become available.
- **Graph provenance is rarely audited.** [8] shows that pooling a knowledge-component graph from the full interaction log before a learner-based train/test split lets held-out information reach a graph-augmented KT model even when the sequence optimizer itself trains only on the training fold. [8]’s protocol audits this channel for *pairwise* prerequisite/similarity edges; no published audit protocol currently covers *hyperedges*.
- **Causal claims about pedagogical intervention are rare.** Most graph-KT models report predictive accuracy only, without attempting to separate a hint or teacher action’s causal effect on subsequent correctness from confounded correlation — motivating an interface for ATE estimation even when public logs do not yet support reporting numbers.

B. Research Questions

This work is motivated by a broader question: *how can a heterogeneous hypergraph – with student, exercise, concept,*

teacher, hint, forum, discussion, and video as node types (Section III) – improve knowledge tracing beyond pairwise student–concept–exercise graphs? We decompose this into five questions answered with evidence below, plus one this paper does not answer:

- **RQ1** (§VI-C): Does a leakage-controlled concept-prerequisite hypergraph yield competitive KT accuracy under a graph-only encoder?
- **RQ2** (§VII-B): Does the resulting model demonstrably rely on the constructed hypergraph, and is that reliance unique relative to reused graph-KT baselines?
- **RQ3** (§VII-C): Does concept-prerequisite hyperedge construction preserve structural fidelity with an expert-curated prerequisite DAG?
- **RQ4** (§VII-D): Can a Critic-gated LLM-agent layer read frozen Tier 1 predictions and produce faithful, useful explanations at pilot scale?
- **RQ5** (§VII-E): Can session+Hint hyperedges and a hint ATE be instantiated on a hint-bearing corpus, and what do propensity diagnostics reveal about identification?

Interactions involving *teacher, forum/discussion, and video* nodes are defined in the schema (Section III) but remain data-blocked on every corpus we use (Section VIII-D): this paper establishes an interface for those entities, not empirical evidence of how they improve KT.

C. Contributions

- **C1.** A hypergraph construction procedure (Section IV-A) that groups [8]’s audited, acyclic prerequisite edges into concept-prerequisite chain hyperedges (*evaluated quantitatively* on XES3G5M/Junyi), plus session hyperedges with Hint members on FoundationalASSIST [9] (Section VII-E); discussion-thread and teacher-intervention types remain data-blocked (Section VIII-D).
- **C2.** A hyperedge-level extension of [8]’s four-diagnostic leakage taxonomy, adding three classes – group-membership, cross-modal, and intervention-timing leakage – that a pairwise audit structurally cannot detect (Section IV-B).
- **C3.** A causal-effect layer (Section IV-D) implementing inverse-propensity ATE estimation; on FoundationalASSIST we report a *limited* observational hint ATE (next-step correctness) with propensity diagnostics (Section VII-E), while XES3G5M still has no synchronized hint treatment for ATE.
- **C4.** Direct reuse of [8]’s published, leakage-controlled baseline results (BKT, DKT, AKT, simpleKT, GKT, GIKT, SKT, DyGKT, DGEKT) as the Tier 1 comparison column under matched protocol and split seeds, avoiding redundant baseline retraining.
- **C5.** A Tier 2 pilot (AgentDG-KT, Section V) showing that a Critic-gated LLM interpretability loop over frozen Tier 1 outputs is tractable at $n=500$ on a consumer GPU, plus a dual-rater Likert study ($n=40$) of faithfulness/usefulness (Table VII) — scoped as pilot evidence, not a full classroom RCT.

- **C6.** Empirical evidence on XES3G5M that (i) a graph-only DH²-KT backbone passes a hypergraph manipulation check ($\Delta\text{AUC}=0.038$ on fold 0) while remaining below reused GKT AUC (mean 0.752 ± 0.002 vs. 0.834 ± 0.001), noting that P0 GKT *also* passes destruction at $p=0.9$ with a larger mean drop, and (ii) the local 7B Tier 2 pilot reaches 0% Critic flag rate after probability-format calibration (Section VI-C).

II. RELATED WORK

A. Graph- and hypergraph-based knowledge tracing

Early graph-augmented KT models propagate information over a single, pairwise concept graph: [2] layer a graph neural network over a knowledge-component graph on top of a sequence model, and [3] let exercises and concepts recursively re-rank one another through a bipartite exercise-concept graph. [7] extends this line to a temporally-evolving graph that updates edges as new interactions arrive rather than freezing the graph after construction. A second line moves from pairwise edges to hyperedges spanning more than two knowledge components at once: [4] proposes a temporal hypergraph memory network, [5] adds a dual-gated mechanism controlling how much a hyperedge’s aggregated signal influences a given prediction, and [6] replaces graph convolution with a hypergraph transformer. A third line adds heterogeneity in what a node *represents* rather than how many nodes an edge spans: [10] augments the concept graph with psychological-factor nodes, and [11] builds a heterogeneous graph directly over learner-question interactions. Across all three lines, every published hyperedge or heterogeneous node we are aware of is still drawn from a closed set of KC-, exercise-, or learner-adjacent entities; none represent a Teacher intervention, Hint content, Forum discussion, or instructional Video as a first-class member inside the same hyperedge as the knowledge components it plausibly affects – the interface gap C1 (Section IV-A) targets, while this paper’s quantitative core evaluates concept-prerequisite chains.

B. LLM agents in education and graph-augmented LLM agents

[12] surveys the growing use of LLM agents across tutoring, feedback generation, and curriculum planning in education, and [13] instantiates a multi-agent LLM pipeline that simulates computerized adaptive testing end to end, with agents standing in for both the test administrator and the examinee. Outside education, [14] surveys how LLM agents can read from and write to an external graph as working memory, and [15] applies a graph-retrieval-augmented-generation pattern to recommend personalized learning paths over a dual knowledge-structure graph. AgentDG-KT (Section V) differs from this pattern in one specific way: none of the above agents are gated by a design in which the graph – and the quantitative model that reads it – is frozen before the agent acts, with the agent limited to producing a natural-language artifact (rationale, hint, session hyperedge) that a separate Critic agent must approve before it is treated as

anything more than a proposal. We adopt this frozen-core-plus-pilot-agent split specifically because a full multi-agent LLM pipeline that re-infers knowledge state end to end over millions of interactions is not tractable on a single consumer GPU (Section V), not because we believe agentic re-inference is undesirable in principle.

C. Leakage-controlled evaluation in knowledge tracing

[8] is the companion paper this work builds on and reuses directly (Section VI-A); we summarize its protocol here so the rest of this paper is self-contained. [8] observes that a knowledge-component graph is often constructed by pooling co-occurrence or similarity statistics over the *entire* interaction log before a learner-based train/test split is applied, which lets held-out information reach a graph-augmented KT model through the graph channel even when the sequence model itself trains only on the training fold. Its audit protocol reports four scalar diagnostics computed on a train-only-constructed graph – an edge-composition-ratio flag (ECR^{flag}), an edge-composition overlap ratio ($\text{ECR}^{\text{overlap}}$), an edge-outcome correlation ($|\rho|$), and a train/test boundary-mismatch rate (TBMR) – alongside a DAG audit with cycle pruning, a cold-start stratification of knowledge components into four frequency strata, a DAG Disruption Rate (DDR) manipulation check under five structured augmentation operators, and a ground-truth cross-validation against an expert-curated prerequisite DAG on Junyi Academy. Across XES3G5M, ASSISTments 2012, and Junyi Academy, [8] reports that train-only graph construction shifts baseline AUC by at most 0.003 relative to full-log construction – evidence that leakage control, in their setting, does not by itself manufacture the reported gains, and a result we rely on directly when interpreting any Tier 1 gain in this paper (Section VI-A). Section IV-B extends this four-diagnostic protocol from pairwise edges to heterogeneous hyperedges.

D. Causal inference in educational data mining

Causal analysis in education has mostly targeted two separate questions: estimating the effect of a known treatment, and discovering an unknown causal structure among skills. [16] addresses the first: given historical course-enrollment and grade data, they estimate the average treatment effect of taking one course before another using matching and regression – directly analogous in spirit to the propensity-score estimator we use for hint/teacher-intervention effects (Section IV-D), though on a coarser course-level treatment rather than a step-level pedagogical action. [17] addresses the second: they learn a latent causal ordering and dependency structure among knowledge-component skills end to end inside a modified DKT model, evaluated on the NeurIPS 2022 Challenge on Causal Insights for Learning Paths in Education. [18] sits between the two, using a causal-relations-constrained student model to make KT predictions more interpretable without a separate treatment-effect estimation step. Our causal-effect layer (Section IV-D) is closest in spirit to [16] – a treatment-effect estimator rather than a causal-discovery method – but applies it at the level of individual hint/teacher-intervention

hyperedges rather than curriculum-level course sequencing, and explicitly reports the propensity-overlap diagnostic that [16] does not discuss, following the identification-assumption caveat in Section VIII-D.

III. PRELIMINARIES

Let $\mathcal{D} = \{(u, t, q, c, y)\}$ be a KT interaction log with learner u , timestamp t , item q , knowledge component c , and correctness label $y \in \{0, 1\}$, following [8]’s notation. A learner-based temporal split $\mathcal{F} = (\mathcal{F}_{\text{train}}, \mathcal{F}_{\text{val}}, \mathcal{F}_{\text{test}})$ assigns each learner to exactly one fold (ratios 0.7/0.1/0.2, three folds with split seeds 42/43/44, matching [8]’s protocol so results remain directly comparable, see Section VI-A).

A. Self-contained summary of the reused P0 protocol

Because [8] is under review at the time of writing, we summarize the reused pieces so this manuscript can be checked without reading the companion paper in full. **Construction.** For each learner-based fold, prerequisite edges E_{pre}^f and similarity edges E_{sim}^f are built from *train-fold interactions only*, then audited for cycles (DAG pruning). **Four pairwise diagnostics** (reported on the train-only graph against held-out learners): edge-composition-ratio flag ECR^{flag} , overlap $\text{ECR}^{\text{overlap}}$, edge–outcome correlation $|\rho|$, and train/test boundary-mismatch rate TBMR. Across XES3G5M, ASSISTments 2012, and Junyi, [8] finds that switching from full-log to train-only graph construction changes baseline AUC by at most 0.003 — evidence that leakage control alone does not manufacture large gains. **Reuse artefacts.** We pin the companion repository to commit 5e3d6a01c32a649f1655ac8712f03359bddfe478 and copy its published baseline CSV tables verbatim for the nine models listed in Section VI-A; config checksum verification for XES3G5M is reported there. Concept-prerequisite hyperedges in this paper are formed by enumerating bounded-length paths on the same audited E_{pre}^f (Section IV-A).

We extend this notation with a *hyperedge* $h = (\kappa, \{(\tau_i, e_i)\}_{i=1}^k, f)$: a kind label $\kappa \in \{\text{concept-prerequisite, session, discussion-thread, teacher-intervention}\}$, a set of k typed member entities (τ_i, e_i) with entity type $\tau_i \in \{\text{student, exercise, concept, teacher, hint, forum, discussion, video}\}$, and a fold identifier f . A heterogeneous hypergraph $\mathcal{H}_f = (\mathcal{V}_f, \mathcal{E}_f^{\text{hyper}})$ for fold f is a set of such hyperedges over the typed node set $\mathcal{V}_f$. **Scope of this paper:** only $\kappa = \text{concept-prerequisite}$ hyperedges (members are concept nodes along audited E_{pre} paths) enter Tier 1 training and the quantitative tables below; the other three kinds are future-proof interfaces plus Tier 2 session write-backs at pilot scale (Section IV-A).

B. Design goal

For each fold f , the goal is a hypergraph $\mathcal{H}_f$ constructed using only $\mathcal{F}_f^{\text{train}}$ evidence, audited for the leakage classes in Table I (Section IV-B), that a downstream Tier 1 encoder consumes to predict y for held-out interactions.




Fig. 1. Two-tier DH^2 -KT architecture. Tier 1 consumes P0’s leakage-audited prerequisite edges as chain hyperedges (solid path; quantitative core on XES3G5M) and optionally FoundationalASSIST session+Hint hyperedges with a limited observational ATE (dashed). After freezing, Tier 2 agents explain and gate those outputs: Diagnostician $\rightarrow$ Critic $\rightarrow$ Tutor, with session hyperedge write-back at pilot scale.

IV. TIER 1: DH^2 -KT

A. Heterogeneous hypergraph construction

Concept-prerequisite hyperedges are built by enumerating simple directed paths of bounded length on [8]’s audited, acyclic E_{pre}^f (Algorithm 1 in code: `build_concept_prerequisite_hyperedges`).

- 1: **Input:** audited, acyclic prerequisite edges E_{pre}^f from [8]’s Algorithm 2; chain length bounds $[\ell_{min}, \ell_{max}]$
- 2: **Output:** concept-prerequisite hyperedge set $\mathcal{E}_f^{cpre}$
- 3: Build directed graph G from E_{pre}^f
- 4: **for** each root r with in-degree 0 in G **do**
- 5: **for** each simple path p from r with $\ell_{min} \leq |p| \leq \ell_{max}$ **do**
- 6: emit hyperedge with members $\{(concept, c) : c \in p\}$
- 7: **end for**
- 8: **end for**

Discussion-thread and teacher-intervention hyperedges remain *data-blocked*: no synchronized Forum or Teacher-action logs ship with the public KT corpora we use. Session hyperedges (Student-Exercise-Concept-Hint) are constructed on FoundationalASSIST [9], which provides `hint_count` synchronized to each interaction (Section VII-E). On the three P0 core benchmarks (XES3G5M, ASSISTments 2012, Junyi), session construction stays blocked because those logs lack Hint/Video columns; ES-KT-24 [19] remains an optional Video-bearing source for future work. Tier 1’s primary predictive core on XES3G5M therefore still uses concept-prerequisite hyperedges; FoundationalASSIST supplies the heterogeneous session+Hint evidence and a secondary predictive fold.

B. Hyperedge-level leakage audit

[8] audits five leakage classes at the level of pairwise edge provenance (structural, edge, target, temporal, cold-start neigh-

TABLE I
HYPEREDGE-LEVEL LEAKAGE TAXONOMY (EXTENDS [8] TABLE 1)

Class	Source	Diagnostic
Group-membership	A hyperedge with ≥ 1 member drawn from a held-out interaction while other members are train-only	<code>compute_group_membership_leakage</code>
Cross-modal	A Forum/Video embedding pooled over content timestamped after the prediction cutoff	<code>compute_cross_modal_leakage</code>
Intervention-timing	A Teacher/Hint hyperedge whose content was effectively chosen using knowledge of the student’s later correctness	<code>compute_intervention_timing_leakage</code>

TABLE II
CONCEPT-PREREQUISITE HYPEREDGE AUDIT ON XES3G5M FOLD 0. PAIRWISE TBMR/ $|\rho|$ ARE SKIPPED WHEN THE PAIRWISE PROJECTION EXCEEDS THE IMPLEMENTATION CAP OF 100,000 PAIRS (11,568,053 PROJECTED PAIRS IN THIS FOLD; SEE AUDIT NOTES).

Metric	Value
Hyperedges	446,305
ECR ^{flag}	0.0
Group-membership leak rate	0.0

bourhood) using four scalar diagnostics: ECR^{flag}, ECR^{overlap}, $|\rho|$, and TBMR. These diagnostics are reused unmodified (Section VI-A) on the pairwise projection of a hyperedge set. Table I adds three classes that a pairwise audit cannot express because they require a *group* of heterogeneous members, not a single edge.

Measured values on XES3G5M fold 0 appear in Table II.

C. Encoding and hypergraph attention

DH²-KT applies two HypergraphConv layers per hyperedge kind on concept embeddings, mean-pools across kinds, and masks concepts not participating in any hyperedge. Sequence encoding uses a *graph-only* interaction term $x_t = \text{proj}(c_t) + \text{proj}(r_t)$ with a learnable graph gain on c_t , deliberately omitting direct exercise embeddings so predictions cannot bypass the hypergraph (Section VI-C). A dual-gated update combines x_t with the previous hidden state; a linear head outputs next-step logits. Training adds a graph-sensitivity auxiliary loss encouraging concept states to depend on hypergraph structure.

D. Causal-effect layer

We estimate the average treatment effect (ATE) of a binary intervention (e.g., hint given) on next-step correctness via inverse-propensity weighting, following the identification assumption that intervention assignment is unconfounded given observed features – a strong assumption on observational KT data that we report as a limitation rather than a proven property (Section VIII-D). Given confounders X , treatment $T \in \{0, 1\}$, and outcome Y , the propensity score $\hat{e}(X) = P(T = 1 | X)$ is estimated by logistic regression, and

$$\widehat{\text{ATE}} = \frac{\sum_i T_i Y_i / \hat{e}(X_i)}{\sum_i T_i / \hat{e}(X_i)} - \frac{\sum_i (1 - T_i) Y_i / (1 - \hat{e}(X_i))}{\sum_i (1 - T_i) / (1 - \hat{e}(X_i))}. \quad (1)$$

No ATE numbers are reported on the primary XES3G5M benchmark (no hint logs). On FoundationalASSIST we estimate an observational IPW ATE of `hint_usedt` on `correctt+1` with pre- t summaries as confounders (Section VII-E); same-step `discrete_score` is *not* used as the outcome because it is mechanically tied to hint requests in that corpus.

V. TIER 2: AGENTDG-KT PILOT

Tier 2 is deliberately scoped as a pilot, not a benchmark at Tier 1’s quantitative rigor (docs/idea-D-plan.md, following the feasibility analysis in the accompanying project discussion: a full multi-agent LLM re-inference pipeline over millions of interactions is not tractable on a single consumer GPU within a realistic project timeline). Three agents operate on Tier 1’s *frozen* output rather than re-inferring knowledge state:

- **Diagnostician Agent** produces a natural-language rationale for a given Tier-1 prediction, explicitly instructed not to override the given probability.
- **Tutor/Hint Agent** proposes the next hint/exercise and writes a new session hyperedge back into the graph (the agentic closed loop).
- **Critic Agent** is a *mandatory* quality gate that flags any Diagnostician output contradicting or ignoring the Tier-1 number it was meant to explain, before anything is treated as ground truth in a case study.

We ran a 500-sample pilot on XES3G5M fold 0 (seed 42) with Tier 1 frozen (checkpoint from graph-only v2 training) and Qwen2.5-7B (Ollama 0.32.6). Table III reports Critic flag rates; v1’s 12.4% flags were predominantly probability *format*

TABLE III
TIER 2 AGENT FAITHFULNESS PILOT (XES3G5M FOLD 0, $n=500$, SEED 42, QWEN2.5-7B VIA OLLAMA). FLAG RATE = FRACTION OF DIAGNOSTICIAN OUTPUTS FLAGGED BY THE MANDATORY CRITIC. V2 APPLIES PROBABILITY-FORMAT CALIBRATION (`CRITIC_CALIBRATION.PY`).

Run	Backend	Flag rate	Mean $P(\text{correct})$
Stub (sanity)	StubLLM	0.000 (0/500)	0.742
Ollama v1	qwen2.5:7b	0.124 (62/500)	0.746
Ollama v2	qwen2.5:7b	0.000 (0/500)	0.753

mismatches (e.g., “69.3%” vs. $P=0.693$), resolved in v2 by `critic_calibration.py` and an updated Critic prompt. The Tutor agent wrote 500 session hyperedges back to the graph, closing the agentic loop at pilot scale.

VI. EXPERIMENTAL SETUP

A. Datasets

Following [8]’s benchmark-role assignment, XES3G5M [20] (18,066 learners, 7,652 items, 865 KCs, 6,413,353 interactions) is the *primary* predictive benchmark: low consecutive-KC-repeat rate, non-saturated headline AUC, and the clearest graph-backbone separation among the corpora [8] studied. All DH²-KT training and the manipulation check reported below use XES3G5M. ASSISTments 2012 is listed as secondary in the P0 protocol (single-skill Q-matrix); we do *not* reuse P0 ASSISTments 2012 GKT AUC as a comparison column. Instead we report a secondary fold on the newer FoundationalASSIST corpus [9] (Section VII-E), which is *not* interchangeable with classic ASSISTments 2012. Junyi Academy supplies the expert prerequisite DAG for *structural* ground-truth overlap only (Section VII-C); full DH²-KT training on Junyi is not run.

B. Baselines reused from P0

The protocol summary in Section III-A underpins reuse. We reuse [8]’s published baseline results (`baselines/p0_reused/baseline_results.csv`, 69 rows, and `baseline_fold_results.csv`, 205 per-fold rows; both copied verbatim from [8]’s `results/tables/`, pinned to commit `5e3d6a01c32a649f1655ac8712f03359bddfe478`, vendored 2026-08-02) for BKT, DKT, AKT, simpleKT, GKT, GIKT, SKT, DyGKT, and DGEKT under train-only, leakage-controlled graph construction, rather than retraining these nine models on a single RTX 3090. Reuse is valid only if our preprocessing exactly matches [8]’s (same hashed `kc_id` space, same Q-matrix). `scripts/01_verify_p0_preprocessing_match.py` checks that each DH²A-KT dataset config points at an unmodified, checksummed copy of [8]’s own config and that the vendored commit is recorded; for XES3G5M this currently confirms the config file matches (`sha256 4f9f35d9...d980b2`) against the vendored commit above. Processed parquet for XES3G5M, Junyi, and ASSISTments 2012 is present under

TABLE IV

TEST AUC ON XES3G5M UNDER MATCHED TRAINING BUDGET (BATCH 4, 10 EPOCHS, MAX SEQ. LEN. 200). P0 FOLD AUCs FROM BASELINES/P0_REUSED (TRAIN-ONLY GRAPHS); DH²-KT v2 USES CHAIN CONCEPT-PREREQUISITE HYPEREDGES AND GRAPH-ONLY INTERACTION. MEAN±SD OVER THREE LEARNER FOLDS.

Fold	DH ² -KT v2	GKT (P0)	simpleKT (P0)
0	0.7501	0.8346	0.8744
1	0.7518	0.8338	0.8747
2	0.7531	0.8326	0.8750
Mean±SD	0.7516±0.0015	0.8336±0.0010	0.8747±0.0003
Δ vs. GKT (mean±SD)		−0.0820±0.0025	

external/p0_leakage_audit/data/processed/. Reviewers may additionally request the companion PDF of [8] as supplementary material while that manuscript is under review.

C. Manipulation check

Following [8]’s Section 4.7 reasoning, we do not interpret any ablation result as evidence of hypergraph benefit unless DH²-KT first passes a manipulation check: AUC must move meaningfully under near-total ($p = 0.90$) destruction of the constructed hyperedge set. A backbone whose AUC does not move under total destruction is *graph-inert*, and its ablation numbers carry no interpretable meaning about graph benefit. On XES3G5M fold 0, DH²-KT v2 passes this gate ($\Delta\text{AUC} = 0.038$; Table V); an earlier graph-inert architecture (pairwise edges + exercise embedding path) failed with $\Delta\text{AUC} \approx 0.0003$, motivating the graph-only redesign.

VII. RESULTS

A. Predictive performance (RQ1)

Table IV reports test AUC on XES3G5M under a training budget matched to [8]’s GKT (batch 4, 10 epochs). DH²-KT v2 mean AUC is 0.7516 versus reused GKT 0.8336 and simpleKT 0.8747. We do not claim SOTA predictive performance; mean±SD over three folds is 0.7516±0.0015 vs. GKT 0.8336±0.0010 (Table IV). The graph-only path is designed so that *our* hypergraph check is interpretable; GKT separately also passes P0’s destruction check (Section VII-B).

B. Manipulation check (RQ2)

Table V and Fig. 2 summarize graph-destruction evidence at the shared P0 anchor ($p=0.9$ node-drop). DH²-KT v2 (fold 0) loses $\Delta\text{AUC} = 0.038$ when frozen weights are re-evaluated on a destroyed hypergraph, confirming hypergraph use under our protocol (an earlier pairwise+exercise-embedding backbone failed this gate, $\Delta\text{AUC} \approx 0.0003$). Reused P0 DDR-downstream results show that **GKT also passes** the same operator/strength on XES3G5M, with a *larger* mean drop (0.0877 ± 0.0009 , $n=9$ fold-seed runs) after retraining on the destroyed pairwise graph. Protocol differences matter (DH²-KT freezes weights; P0 GKT retrain), but either way graph-reliance is *not* unique to DH²-KT. simpleKT has no concept-graph adjacency, so ΔAUC under graph destruction is undefined (N/A). Fig. 2 draws an *illustrative* horizontal line at

TABLE V

MANIPULATION CHECK AT NODE-DROP $p=0.9$ ON XES3G5M. DH²-KT: TRAIN ON CLEAN HYPEREDGES, RE-EVALUATE FROZEN WEIGHTS ON DESTROYED HYPEREDGES (FOLD 0). GKT: P0 DDR-DOWNSTREAM MEAN ΔAUC OVER $n=9$ FOLD-SEED RUNS ([8]; RETRAIN ON DESTROYED PAIRWISE GRAPH). SIMPLEKT HAS NO CONCEPT-GRAPH INPUT, SO GRAPH DESTRUCTION IS UNDEFINED.

Model	ΔAUC ($p=0.9$)	DDR	Pass?
DH ² -KT v2 (fold 0)	0.0380	0.990	Yes
GKT (P0, mean±SD)	0.0877 ± 0.0009	0.990	Yes
simpleKT (P0)	N/A	—	N/A




Fig. 2. Manipulation-check ΔAUC on XES3G5M fold 0 under $p=0.9$ node-drop. An early pairwise+exercise-embedding backbone is graph-inert ($\Delta\text{AUC} \approx 0.0003$); the chain+graph-only redesign passes ($\Delta\text{AUC} = 0.038$). The dashed line at 0.003 is an *illustrative* floor ($\approx 10 \times$ the graph-inert noise), not a formal statistical cutoff from [8].

$\Delta\text{AUC} = 0.003$: roughly an order of magnitude above the observed graph-inert noise floor ($\approx 3 \times 10^{-4}$), not a formal hypothesis test from [8]. We treat a backbone as graph-reliant when its destruction ΔAUC clearly exceeds that floor (here 0.038 for DH²-KT; ≈ 0.088 for GKT).

C. Ground-truth structural overlap (Junyi)

Table VI compares hyperedge-chain overlap with P0’s pairwise E_{pre} against the Junyi expert DAG at $K=|\text{expert}|$. Edge F1 is 0.165 (hyperedge chain) vs. 0.183 (E_{pre}) on fold 0; folds 1–2 are analogous. The chain representation slightly *underlaps* pairwise E_{pre} : collapsing directed paths into multi-node hyperedges, then projecting back to pairs for overlap scoring, can dilute expert-edge recall relative to scoring E_{pre} directly. We treat this as structural validation of construction fidelity, *not* as evidence that chain hyperedges improve predictive KT on Junyi — full DH²-KT training on Junyi (25M interactions) was not run.

D. Tier 2 faithfulness pilot

See Table III and Section V. The Tutor wrote 500 session hyperedges at pilot scale. From the calibrated Ollama v2 run we drew a fixed 40-sample pack (seed 42) for two independent raters (A1, B1) on Likert faithfulness and usefulness (1–5).

TABLE VI

GROUND-TRUTH CROSS-VALIDATION ON JUNYI ACADEMY (FOLD 0, TOP- $K=|\text{EXPERT}|=1131$). HYPEREDGE CHAIN SLIGHTLY UNDERLAPS PAIRWISE E_{pre} BUT REMAINS THE SAME ORDER OF MAGNITUDE.

Representation	Edge F1	Edge prec.	Edge rec.
P0 E_{pre}	0.183	0.183	0.183
Hyperedge chain	0.165	0.165	0.165

TABLE VII

HUMAN RATING OF TIER 2 OUTPUTS (XES3G5M FOLD 0 OLLAMA V2 PILOT SUBSAMPLE, $n=40$ EXPLANATIONS, 2 INDEPENDENT RATERS, LIKERT 1–5). AGREEMENT IS PAIRWISE SPEARMAN ρ AND FRACTION OF SCORES WITHIN ± 1 .

Scale	Mean $\pm$ SD	Agreement (A1 vs. B1)
Faithfulness	4.01 ± 1.07	$\rho=0.71$; within ± 1 : 92.5%
Usefulness	3.53 ± 0.71	$\rho=0.58$; within ± 1 : 97.5%

Pooled means are 4.01 ± 1.07 (faithfulness) and 3.53 ± 0.71 (usefulness); pairwise Spearman $\rho=0.71$ / 0.58 and within- ± 1 rates of 92.5% / 97.5% (Table VII). A third rater was optional and not required for this pilot report.

E. Secondary corpus: FoundationalASSIST (session+Hint)

FoundationalASSIST [9] provides $\approx 1.74\text{M}$ skill-linked interactions with synchronized `hint_count` (hint-use rate 5.2%). Using the same learner-temporal protocol as P0 (0.7/0.1/0.2, fold 0), we infer train-only E_{pre} (280 edges after cycle pruning), build 5,840 concept-prerequisite chain hyperedges, and construct 308,670 session hyperedges on train users (11.6% with a Hint member; $\text{ECR}^{\text{flag}}=0$, group-membership leak=0). DH²-KT fold 0 AUC is 0.720 with concept-prerequisite hyperedges alone ($n=298,500$ predictions). Adding up to 20k multi-concept session hyperedges (Hint-preferring selection; concept co-occurrence projection) raises AUC to 0.722 ($\Delta+0.002$) — a small positive shift under the same local budget, not a matched GKT comparison. Using the causal layer on train transitions, the IPW ATE of hint use on *next-step* correctness is -0.207 (naive mean difference -0.255 ; trimmed IPW -0.211 ; bootstrap 95% CI $[-0.230, -0.187]$ on 50k-row resamples; $n_{\text{treated}}=36,151$). Propensity scores hit the clip floor ($[0.05, 0.219]$), indicating weak overlap: struggling learners disproportionately request hints, so the negative association is *not* interpretable as a proven causal penalty of giving hints. Table VIII summarizes.

VIII. DISCUSSION

A. Threats to validity

Our XES3G5M comparison uses [8]’s reused baselines under matched batch/epoch budget (Table IV); preprocessing configs are pinned to P0’s vendored commit. DH²-KT mean AUC is 0.7516 ± 0.0015 vs. GKT 0.8336 ± 0.0010 (paired $\Delta=-0.082 \pm 0.003$ over three folds). Both DH²-KT and GKT pass strong graph-destruction checks; the AUC gap

TABLE VIII

FOUNDATIONALASSIST SECONDARY CORPUS (FOLD 0). SESSION HYPEREDGES USE A 30-MINUTE GAP; HINT MEMBERS APPEAR WHEN `HINT_COUNT > 0`. PREDICTIVE AUC IS OBSERVATIONAL (NO P0-REUSED GKT/SIMPLEKT COLUMN). HINT ATE IS IPW OF `HINT_USEDt` ON `CORRECTt+1` UNDER PRE- t CONFOUNDERS; PROPENSITY MASS HITS THE CLIP FLOOR — TREAT AS LIMITED OBSERVATIONAL EVIDENCE.

Quantity	Value
Interactions (canonical)	1,740,987
Learners / items / KCs	5,000 / 3,359 / 209
Hint-use rate	0.052
Train-only E_{pre} edges (after cycle prune)	280
Concept-prerequisite chain hyperedges	5,840
Session hyperedges (train users)	308,670
with ≥ 1 Hint member	35,752 (11.6%)
Session hyperedge ECR^{flag} / group leak	0 / 0
DH ² -KT fold 0 AUC (concept-prereq)	0.720
+ session HE (20k, Hint-prefer)	0.722 ($\Delta+0.002$)
n next-step predictions	298,500
Hint ATE (IPW) / naive mean diff	-0.207 / -0.255
trimmed IPW (propensity 10–90%)	-0.211
bootstrap 95% CI (subsample)	$[-0.230, -0.187]$
n treated / control (train transitions)	36,151 / 660,321
propensity [min, max]	$[0.05, 0.219]$

therefore cannot be read as “DH²-KT uniquely buys graph-reliance.” Section VIII-B repositions value around audit, heterogeneous/causal interfaces, and Tier 2 rather than exclusive graph sensitivity. Primary claims remain scoped to concept-prerequisite hyperedges on XES3G5M; FoundationalASSIST is secondary (Section VII-E). Passing destruction confirms graph use for a backbone, not that a specific hyperedge kind drives predictive gain.

B. What justifies DH²-KT if graph-reliance is not unique?

Because GKT also passes destruction while beating DH²-KT on clean AUC (Table V), the remaining case for DH²-KT is what reused GKT does *not* provide: (i) hyperedge-level leakage diagnostics beyond pairwise E_{pre} (Table I); (ii) a shared schema for heterogeneous members (session/Hint/teacher/forum/ video), exercised for session+Hint on FoundationalASSIST while discussion-thread and teacher types stay data-blocked; (iii) an inverse-propensity ATE layer for pedagogical interventions (Section IV-D; limited observational hint ATE on FoundationalASSIST); and (iv) a Critic-gated explanation loop over frozen $P(\text{correct})$ (Section V). Protocol note: DH²-KT ΔAUC is frozen-weight re-evaluation; P0 GKT ΔAUC retrains on the destroyed pairwise graph — complementary, not apple-to-apple. **Prefer high-AUC baselines** when ranking alone is the product goal; **prefer DH²-KT** when stakeholders need higher-order auditability, an extensible heterogeneous/causal interface, and/or Critic-gated explanations, accepting the ≈ 8 point AUC gap versus GKT (-0.082 ± 0.003 ; Table IV) as the cost of our graph-only path — not a claim that GKT is graph-inert.

C. Tier 2 faithfulness

Tier 2’s 12.4% v1 flag rate was largely a Critic prompt artifact (% vs. decimal); calibration reduced it to 0% on 500 samples without changing Tier 1 numbers. Dual-rater scores

(Table VII) place faithfulness near “aligned” ($\approx 4/5$) and usefulness in the “somewhat–mostly useful” band ($\approx 3.5/5$), with high within- ± 1 agreement; usefulness Spearman is moderate, reflecting more subjective teacher-facing judgments.

D. Limitations

- Discussion-thread and teacher-intervention hyperedges remain data-blocked; session+Hint hyperedges are demonstrated on FoundationalASSIST (concept projection; +0.002 AUC vs. concept-only) but are not the primary XES3G5M training graph.
- The causal-effect layer’s identification assumption (no unmeasured confounding given observed features) is stated, not proven; the FoundationalASSIST hint ATE shows weak propensity overlap and should not be read as a recommendation against hints.
- The hyperedge leakage audit’s pairwise-projection approximation (Section IV-B) weights all projected pairs equally; a hyperedge-native statistic is future work.
- Tier 2 pilot: one fold, one local 7B LLM; Critic sensitivity to response format required calibration; human ratings use $n=40$ and two raters (no third rater; not a classroom deployment study).
- Reproduction environment: Windows workstation, single RTX 3090; 42/42 unit tests pass in the released codebase.

IX. CONCLUSION AND FUTURE WORK

We presented DH²A-KT, a two-tier framework combining leakage-controlled hypergraph knowledge tracing (DH²-KT) with a frozen-core LLM agent pilot. On XES3G5M, graph-only DH²-KT and reused GKT both pass strong destruction checks; DH²-KT remains below GKT AUC (0.752 ± 0.002 vs. 0.834 ± 0.001) under a graph-only interaction path. Quantitative claims are scoped to concept-prerequisite chain hyperedges. On Junyi, those hyperedges show structural overlap with the expert DAG (edge-F1 ≈ 0.17) without full-model training. Tier 2 demonstrates a Critic-gated loop at 500-sample scale with 0% flag rate after calibration and dual-rater Likert means of $\approx 4.0 / 3.5$ for faithfulness / usefulness. On FoundationalASSIST we add session+Hint hyperedges, a secondary fold-0 AUC, and a limited observational hint ATE with propensity diagnostics. Future work includes richer confounders / sensitivity analysis for ATE, a third rater or larger human sample, ES-KT-24 Video wiring, and hyperedge-native leakage statistics.

ACKNOWLEDGMENT

All Tier 1 experiments in this paper are designed to run on a single consumer-grade NVIDIA RTX 3090 (24GB), matching the compute budget already validated by [8] on the same primary benchmarks. FoundationalASSIST [9] was accessed under the Hugging Face dataset’s Responsible Use Guidelines for this study.

REFERENCES

- [1] C. Piech, J. Bassen, J. Huang, S. Ganguli, M. Sahami, L. J. Guibas, and J. Sohl-Dickstein, “Deep knowledge tracing,” in *Advances in Neural Information Processing Systems*, vol. 28, 2015, pp. 505–513. [Online]. Available: <https://doi.org/10.48550/arXiv.1506.05908>
- [2] H. Nakagawa, Y. Iwasawa, and Y. Matsuo, “Graph-based knowledge tracing: Modeling student proficiency using graph neural networks,” in *IEEE/WIC/ACM International Conference on Web Intelligence (WI)*, 2019, pp. 156–163. [Online]. Available: <https://doi.org/10.1145/3350546.3352513>
- [3] Y. Yang, J. Shen, Y. Qu, Y. Liu, K. Wang, Y. Zhu, W. Zhang, and Y. Yu, “GIKT: A graph-based interaction model for knowledge tracing,” in *Machine Learning and Knowledge Discovery in Databases (ECML-PKDD)*, 2020, pp. 299–315. [Online]. Available: https://doi.org/10.1007/978-3-030-67658-2_18
- [4] M. Mohammadi, K. Berahmand, S. Sadiq, and H. Khosravi, “Knowledge tracing with a temporal hypergraph memory network,” in *Artificial Intelligence in Education. Posters and Late Breaking Results, Workshops and Tutorials, Industry and Innovation Tracks, Practitioners, Doctoral Consortium, Blue Sky, and WideAIED*, ser. Communications in Computer and Information Science, vol. 2592. Springer, 2025, pp. 77–85. [Online]. Available: https://doi.org/10.1007/978-3-031-99267-4_10
- [5] F. Ma, C. Zhu, and P. Lei, “Hypergraph-driven high-order knowledge tracing with a dual-gated dynamic mechanism,” *Applied Sciences*, vol. 15, no. 15, p. 8617, 2025. [Online]. Available: <https://doi.org/10.3390/app15158617>
- [6] X. Guo, X. Chang, X. Zhang, K. Liang, Y. Shang, and Y. Zhang, “HTKT: Knowledge tracing based on hypergraph transformer,” in *IEEE International Symposium on Parallel and Distributed Processing with Applications (ISPA)*. IEEE, 2024, pp. 1472–1477. [Online]. Available: <https://doi.org/10.1109/ISPA63168.2024.00199>
- [7] K. Cheng, L. Peng, P. Wang, J. Ye, L. Sun, and B. Du, “DyGKT: Dynamic graph learning for knowledge tracing,” in *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)*, 2024, pp. 409–420. [Online]. Available: <https://doi.org/10.1145/3637528.3671773>
- [8] D. M. Tuan, K.-T. Nguyen, D. N. Tien, Q. K. Ngo, V.-H. Nguyen, and L. H. Son, “Leakage-controlled concept graph construction and cold-start diagnostic protocol for knowledge tracing,” 2026, submitted to Applied Intelligence (APIN); code and artefacts at <https://github.com/edu-risk-lab/leakage-controlled-kt-audit>.
- [9] ASSISTments Foundation, “FoundationalASSIST: An educational dataset for foundational knowledge tracing and pedagogical grounding of LLMs,” arXiv:2602.00070, 2026, hugging Face dataset: <https://huggingface.co/datasets/ASSISTments/FoundationalASSIST>. [Online]. Available: <https://doi.org/10.48550/arXiv.2602.00070>
- [10] Z. Wang, W. Wu, C. Zeng, H. Luo, and J. Sun, “Psychological factors enhanced heterogeneous learning interactive graph knowledge tracing for understanding the learning process,” *Frontiers in Psychology*, vol. 15, p. 1359199, 2024. [Online]. Available: <https://doi.org/10.3389/fpsyg.2024.1359199>
- [11] Q. Ni, T. Wei, J. Zhao, L. He, and C. Zheng, “HHSKT: A learner–question interactions based heterogeneous graph neural network model for knowledge tracing,” *Expert Systems with Applications*, vol. 215, p. 119334, 2023. [Online]. Available: <https://doi.org/10.1016/j.eswa.2022.119334>
- [12] Z. Chu, S. Wang, J. Xie, T. Zhu, Y. Yan, J. Ye, A. Zhong, X. Hu, J. Liang, P. S. Yu, and Q. Wen, “LLM agents for education: Advances and applications,” in *Findings of the Association for Computational Linguistics: EMNLP 2025*, 2025, pp. 13 832–13 870. [Online]. Available: <https://doi.org/10.18653/v1/2025.findings-emnlp.743>
- [13] W. Zhou, H. Ma, X. Yu, C. Wang, S. Yang, and X. Zhang, “AgentCAT: Simulating computerized adaptive testing via multi-agent large language models,” arXiv:2606.21832, 2026. [Online]. Available: <https://doi.org/10.48550/arXiv.2606.21832>
- [14] Y. Liu, G. Zhang, K. Wang, S. Li, and S. Pan, “Graph-augmented large language model agents: Current progress and future prospects,” *IEEE Intelligent Systems*, 2025. [Online]. Available: <https://doi.org/10.1109/MIS.2025.3642667>
- [15] X. Cheng, Z. Zhang, J. Wang, L. Fang, C. He, Q. Guan, S. Pan, and W. Luo, “GraphRAG-induced dual knowledge structure graphs for personalized learning path recommendation,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 40, no. 17, 2026. [Online]. Available: <https://doi.org/10.1609/aaai.v40i17.38479>

- [16] P. Kaur, A. Polyzou, and G. Karypis, “Causal inference in higher education: Building better curriculums,” in *Proceedings of the Sixth (2019) ACM Conference on Learning @ Scale*, 2019, pp. 1–4. [Online]. Available: <https://doi.org/10.1145/3330430.3333663>
- [17] N. A. Kumar, W. Feng, J. Lee, H. McNichols, A. Ghosh, and A. Lan, “A conceptual model for end-to-end causal discovery in knowledge tracing,” in *Proceedings of the 16th International Conference on Educational Data Mining (EDM 2023)*, 2023, pp. 413–418. [Online]. Available: <https://doi.org/10.5281/zenodo.8115727>
- [18] S. Minn, J.-J. Vie, K. Takeuchi, H. Kashima, and F. Zhu, “Interpretable knowledge tracing: Simple and efficient student modeling with causal relations,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 36, no. 11, 2022, pp. 12 810–12 818. [Online]. Available: <https://doi.org/10.1609/aaai.v36i11.21560>
- [19] D. Kim, U. Lee, S. Lee, J. Bae, T. Ahn, J. Park, G. Lee, and H. Kim, “ES-KT-24: A multimodal knowledge tracing benchmark dataset with educational game playing video and synthetic text generation,” in *Artificial Intelligence in Education*, ser. Lecture Notes in Computer Science, vol. 15882. Springer, 2025, pp. 281–295. [Online]. Available: https://doi.org/10.1007/978-3-031-98284-2_20
- [20] Z. Liu, Q. Liu, T. Guo, J. Chen, S. Huang, X. Zhao, J. Tang, W. Luo, and J. Weng, “XES3G5M: A knowledge tracing benchmark dataset with auxiliary information,” in *Advances in Neural Information Processing Systems 36 (NeurIPS 2023), Datasets and Benchmarks Track*, 2023, pp. 32 958–32 970. [Online]. Available: <https://doi.org/10.52202/075280-1429>